# Stoa Guide (Tamil Singapore track)

## 1. Welcome > Setup

Skip this if you already worked on Octahedron i18n Comparisons; it is the same VPN and Hatch login.

1. Install OpenVPN Connect: go to openvpn.net/client, download OpenVPN Connect for your system and install it.
2. Add the Handshake server URL: open the app, choose to connect with a Server URL, type https://openvpn.handshakecommunity.ai/ and click Continue.
3. Sign in and import your profile: sign in to the portal with your VPN username and password (it looks like yourid-yourproject). Download your profile file; if OpenVPN Connect does not load it, open the file with the app. Use Google Authenticator for 2-step verification.
4. Connect: select your profile and click Connect. You are connected when the app shows Securely Connected and a green dot.
5. Sign in to Hatch: keep the VPN on and sign in to Hatch with the link and login Handshake gave you. If Hatch does not load, check the VPN still shows Securely Connected.
6. Use your alias: session names start with your alias (example 27b6b2b03e31). Use it exactly as given, in lowercase. When you create the Hatch session, check the persona matches your universe in Your universe and persona.

## 2. What changed from Octahedron i18n Comparisons

Most of Stoa is familiar: same Hatch protocol, same eight dimensions, same scale, no ties, absolute quality, and an English rationale with quoted words.

Short version: you no longer write the prompt; it comes with the task. The form drafts a suggested verdict for you to check, you enter the verdict in Hatch first and then copy it to the form, and you have 20 minutes. A rejected task is final.

What changed (Octahedron i18n Comparisons -> Stoa):

- Prompt: you explored Apps and wrote the prompt -> Given on the task, in your language, with an English gloss and the records it is grounded in. Small wording changes only; never change the asks, names, dates, amounts, language or script.
- Category: you picked one -> Comes with the task.
- Follow-up turns: not allowed -> Not allowed. Send only the task prompt, then evaluate the two replies. Do not send another message to either arm.
- Universe: you picked one of your locale's universes -> One fixed universe for all your Stoa tasks, shown on the task.
- First step: explore Apps and write the prompt -> Upload a screenshot of the universe record the prompt is based on.
- AI help: no AI of any kind -> The form generates a suggested verdict (scores, preference, absolute quality, rationale notes, or SKIP). Check it, own it, rewrite it in your own words. No outside AI tools.
- Order: fill HAI form and clear checks first, then Submit preference and End comparison -> Enter the verdict in Hatch first, Submit preference and End comparison, then record the same values on the form.
- Rationale: English, more than five sentences, quoted words explained -> Same, written by you from the suggestion's notes. Also: no bullet points, no score list, no em dashes.
- Form fields: universe name, category, apps utilized, transcripts -> Four session confirmations, session name and ID, prompt as sent, Response A and B, Trace A and B, then the verdict.
- Screenshots: response, scoring panel, confirmed status -> Universe record at the start, and the submitted Hatch scoring panel at the end.
- Paid time: 65 minutes -> 20 minutes per task, Hatch plus the form.

- Returned tasks: came back with feedback -> Rejections are final. Check everything before you submit.
- Help: project channels -> #stoa-tasking.

Stays the same:

- Same VPN and Hatch login.
- Mode GAIA PREFERENCE and protocol Hatch RM Style Task - Multilingual v1.
- Session names alias_task_N, abandoned sessions alias_abandonedtask_N, session ID copied exactly.
- Click Start annotation before sending the prompt.
- The same eight dimensions on the same 1 to 5 A-versus-B scale.
- Preference is A or B, never a tie, even though Hatch shows a Tie button.
- Absolute quality of the preferred arm: Perfect, Good or Bad.
- Rationale in English, more than five sentences, decisive signal first, original words quoted and explained.
- Trace A and Trace B from each arm's own Copy trace button.
- Notes are required; write "None" if there is nothing.
- Never submit a broken comparison; retry, or abandon and recreate.

Unlearn these (Good vs Bad):

- Good: use the given prompt, with small wording tweaks at most. Bad: explore Apps and write your own prompt.
- Good: send only the task prompt, then evaluate the two replies. Bad: send another message to either arm.
- Good: enter the verdict in Hatch first, then copy it to the form. Bad: fill the form first and submit in Hatch last.
- Good: check the suggested verdict, then write the rationale in your own sentences. Bad: avoid the suggestion entirely, or paste its notes as your rationale.
- Good: write the rationale as prose, with no bullets and no em dashes. Bad: lay out one point per line.
- Good: get it right the first time within 20 minutes. Bad: count on a returned task and a 65-minute budget.

## 3. Your universe and persona

Your task shows your universe id. Hatch shows a persona name. They are the same person. Use that persona's name when you create your Hatch session and when you judge whether a reply sounds like them.

Tamil (Singapore) (our track):

- omni-ta-sg-01 = Keerthana Palanisamy
- omni-ta-sg-02 = Ramesh Chandran
- omni-ta-sg-03 = Murugan Shanmugam
- omni-ta-sg-04 = Prakash Ganesan
- omni-ta-sg-05 = Shanthi Anand
- omni-ta-sg-06 = Janani Shanmugam
- omni-ta-sg-07 = Divya Balan
- omni-ta-sg-08 = Ashwin Chandran
- omni-ta-sg-09 = Tamilselvi Krishnan

Other locales:

- French (Quebec): 01 Alexandre Fournier, 02 Simon St-Pierre, 03 Louis Turcotte, 04 Sylvain Ouellet, 05 Charlotte Gauthier, 06 Stephanie Bouchard, 07 Emma Poirier, 08 Denis Comeau.
- Spanish (Mexico): 01 Guadalupe Estrada Ruiz, 02 Antonia Hernandez Castillo, 03 Sebastian Medina Garcia, 04 Hugo Lopez Silva, 05 Rocio Rodriguez Campos, 06 Esperanza Campos Gomez, 07 Francisco Vargas Gutierrez, 08 Adrian Castro Navarro, 09 Brenda Vargas Torres.
- Portuguese (Brazil): 01 Andre Flores, 02 Rosangela Gomes, 03 Fabiana Bezerra, 04 Mariana Lopes, 05 Jorge Sarkis, 06 Marina Gallo, 07 Jamile Tannous, 08 Renata Nasser.

- Japanese (omni-ja-jp-01 to 08): 8 personas (names in Japanese script on the site).
- Chinese (Singapore): 01 Chng Li Ting, 02 Toh Kai Jie, 03 Chong Hui Ling, 04 Heng Kai Wen, 05 Wong Wan Ling, 06 Pang Hao Ran, 07 Tay Li Xuan, 08 Tan Jia Hui, 09 Loh Ming Hui.
- Malay (Singapore): 01 Nabila Hassan, 02 Iskandar Ismail, 03 Adam Hassan, 04 Yasmin Basri, 05 Irfan Basri, 06 Zulkifli Daud, 07 Iskandar Ahmad, 08 Sofia Ismail, 09 Haris Hassan.
- Bengali (Singapore): 01 Sayan Pramanik, 02 Kaushik Kundu, 03 Madhumita Sen, 04 Sujoy Guha, 05 Rudra Mukherjee, 06 Srabani Majhi, 07 Kamrul Hossain, 08 Anima Guha, 09 Ankita Bandyopadhyay.
- Hindi (Singapore): 01 Arjun Mehra, 02 Anjali Jain, 03 Aditya Jain, 04 Aditi Khanna, 05 Kriti Saxena, 06 Rahul Sharma, 07 Harjit Gill, 08 Jyoti Yadav, 09 Vikram Kapoor.
- Hindi (India): 01 Pavneet Bajwa, 02 Kamla Dubey, 03 Priya Singh, 04 Kiara Malhotra, 05 Aakash Bansal, 06 Faisal Ansari, 07 Dinesh Yadav, 08 Amit Upadhyay, 09 Shaista Hashmi.
- Bengali (India): 01 Aparajita Banerjee, 02 Jyotirmoy Guha, 03 Sohail Laskar, 04 Ila Sarkar, 05 Mira Sarkar, 06 Chaya Pramanik, 07 Anwar Mistri, 08 Chandan Karmakar, 09 Ruksana Khan.
- English (India): 01 Harpreet Singh Pannu, 02 Christopher Rozario, 03 Aravind Chettiar, 04 Anjali George, 05 Aaliya Khan, 06 Aishwarya Kamath, 07 Meena Iyengar, 08 Rahul Joseph, 09 Kiran Shetty.

If your universe is not listed, or the persona in Hatch does not match, stop and ask in the project Slack channel before you task.

## 4. Task workflow

1. Claim the task: claim the Stoa task on the Handshake AI platform and keep it open. The open task is your paid clock.
2. Read your prompt: read the prompt, the English gloss, the category and the records it is grounded in.
3. Add the universe screenshot: connect to the VPN, open Hatch and find the record the prompt is based on (the email, chat, event or file) in your universe, then upload a screenshot of it. This shows you know the context behind the prompt.
4. Create the Hatch session: Mode GAIA PREFERENCE, your universe, Protocol Hatch RM Style Task - Multilingual v1. Wait until the session shows RUNNING before you start.
5. Rename the session: alias_task_N (your alias, then _task_, then the task number). Example: 27b6b2b03e31_task_12.
6. Confirm the session on the form: tick the four confirmations (mode, universe, protocol, session name). A session on the wrong universe or protocol is not usable; abandon it and create a new one.
7. Copy the session name and session ID into the form. The session ID looks like 2e6c6e89-82af-44b4-9870-186ac9e29bb9. Copy it character for character.
8. Start the annotation and send the prompt: click Start annotation, then paste your prompt. Turns sent before the annotation starts are not captured.
9. Check both arms finished: no missing reply, no error, no cut-off text. Check this before you evaluate anything.
10. Paste what you ran: the prompt as sent, Response A, Response B, Trace A and Trace B (Copy trace on each arm). Paste exactly as it is in Hatch. Do not edit, trim or translate.
11. Read the suggested verdict: the form generates scores, a preference, an absolute quality and rationale notes. Read all of it. Check that every quoted phrase really appears in the arm it is attributed to.
12. Enter the verdict in Hatch: set the eight scores, the preference and the absolute quality. Write the rationale in your own full sentences, in English. Click Submit preference, then End comparison.
13. Record the same values on the form: preference, absolute quality, the eight scores and the rationale, word for word as submitted in Hatch.
14. Upload the panel screenshot: the submitted Hatch scoring panel showing the scores, preference, absolute quality and rationale. Use up to three if it does not fit in one. Show the session ID if you can.
15. Write Notes and submit: note any disagreement with the suggestion, apps the arms could not reach, or Hatch errors. Write "None" if there is nothing. Confirm and submit.

## 5. Checking what the arms did

The traces show what each arm actually did. Read them, not only the chat.

Five mistakes reviewers watch for:

1. Trace-contradicted claims: the reply says it sent or created something the trace does not show.
2. Invented details praised as personalisation: names, dates or amounts that appear in no tool result are a Correctness fault, never a Style win.
3. Unrequested writes: an arm that sent instead of drafting, added attendees or deleted something is not rewarded for initiative.
4. Completion claims taken at face value: "Done, I created the doc" with no write in the trace is not done.
5. A winner that does not match the reasoning: check that the selected winner agrees with the explanation you wrote.

What each app can do (an arm claiming something the app cannot do loses on Correctness):

- Gmail: read, real drafts, send. A requested draft must exist as a draft. Sending when asked to draft is a serious failure.
- Google Calendar: create and change events. Check date, year, time zone and attendees.
- Google Contacts: read only.
- WhatsApp: read and search; send only to an existing chat. Cannot create groups.
- Outlook Mail: read and send. Cannot create drafts or move mail.
- Google Drive, Docs, Sheets, Notion: create and change. Describing a document is not creating one.
- Oura: read only. Each sleep row is dated the morning the person woke up.

Incomplete traces: a trace can miss work done by sub-steps. If an arm reports specific data and its trace shows no matching call, check it in Hatch before calling it invented.

Hatch dates: Hatch stamps new actions with the real clock, not the persona's pretend date. Judge the record that was made.

## 6. Golden examples

These rationales show the evidence and structure reviewers look for. Your own rationale must be more than five full sentences.

Good (French Quebec, B slightly better, Good): "Both arms read the same two emails and report the facts correctly, so Correctness is a Tie; neither noticed Alexandre had already replied to Isabelle, hence Good, not Perfect. Helpfulness and Completeness are also a Tie, since both answer his question with the same two emails. He had two worries, looking pushy and missing something; B's closing 'Pas besoin de le relancer, t'es correct' ('no need to chase him, you're fine') answers both in his own casual Quebecois, A only the first. That line is why Tone and Warmth go to B. Style: B says 't'es correct', which is what we'd say here, while A's 'jeter un oeil' ('take a look') is neutral international French; B's 'en lice' ('in the running') is a bit formal, hence only slightly. Coherence goes to A: 'du nouveau d'aujourd'hui' is clunky."

Bad: "B is slightly better because it sounds more natural and has a better tone. Style 4, Tone 4, Warmth 4, Correctness 3. Overall, both responses are good but B is warmer."

Why bad: no quoted words, restates the scores, and gives no reason a reviewer can check.

Good (Hindi India, A significantly better, Bad; separates the preferred arm from its absolute quality): "Both arms failed: neither opened Gmail, Docs, Notion, Calendar or Drive, both asked Pavneet for figures already in her '28k hisaab rough' doc, and neither set the RD reminder; that is the Bad. A still delivered two of her four asks, B none: A created the 'October Budget Plan' note it promised, B promised a plan and did nothing. A's note wrongly says a missed Post Office RD instalment hurts her CIBIL score, so Correctness is only slightly A. Register: Pavneet wrote in friendly 'tum'; A keeps it, B slips into intimate 'tu', unearned for an assistant, and brushes off a user who said she was very worried with 'don't stress'. That is why Style, Tone and Warmth go to A. Coherence goes to B:

A is three stacked messages repeating 'I understand', B one tidy reply."

Bad: "A is better because it created the budget note. B did not. A is Perfect."

Why bad: rates the preferred arm Perfect while it missed two of four asks and made a factual error. Preferred does not mean good.

Session name. Good: 27b6b2b03e31_task_12. Bad: 27B6B2B03E31 Task 12, or 2e6c6e89-82af-44b4-9870-186ac9e29bb9.

Why: uppercase and spaces break the format, and the second is the session ID, not the name.

Broken comparison. Good: Response B failed, so the session was abandoned, a new session was created and the task was run from the beginning. Bad: submitting with "Response B failed" in Notes and scores filled in anyway.

Why: Notes cannot repair a missing arm. A broken comparison must never be submitted.

## 7. Common errors

- Pasting the suggestion's bullets as your rationale. Good: rewrite the points in your own full sentences. Bad: copy the bullet list into Hatch.
- No quoted evidence. Good: "B's 't'es correct' ('you're fine') is what a friend here says". Bad: "B sounds more natural".
- Restating the scores. Good: say what each arm did where something stood out. Bad: "Tone: B slightly better because B had the better tone".
- Defaulting to 3. Good: name a real difference and score 2 or 4, or name the shared trait behind the 3. Bad: score 3 because the arms feel close.
- Choosing a tie preference. Good: pick A or B, slightly if the gap is small. Bad: click Tie in Hatch.
- Rewarding length. Good: credit extra length only for useful content or action. Bad: prefer the longer reply because it says more.
- Form values that do not match Hatch. Good: copy the exact values and rationale you submitted in Hatch. Bad: adjust scores on the form after submitting in Hatch.
- Wrong session. Good: GAIA PREFERENCE, your universe, Hatch RM Style Task - Multilingual v1. Bad: a session on another universe, another protocol or a persona from a different universe.
- Mixing up session name and session ID. Good: name alias_task_N, ID the UUID, each in its own field. Bad: the UUID in the session-name field.

## 8. When Hatch breaks

One rule governs everything here: never submit an incomplete comparison.

A comparison is incomplete if:

- Response A or Response B is missing
- Either arm shows only an error such as "Something went wrong"
- Either arm is cut off
- One arm did not receive all the user turns
- You copied text from one arm into the other
- The comparison ended before you submitted the preference

Recovery steps:

1. Retry the failed arm. A single failed reply can often be rerun.
2. Abandon the session. If it cannot be recovered, rename it alias_abandonedtask_N, for example 27b6b2b03e31_abandonedtask_12.
3. Create a new session: same mode, same universe, same protocol, name alias_task_N.

4. Run the task from the beginning: start the annotation first, then send the prompt. Do not paste output from the broken session.
5. Report repeated problems: post the pattern in #stoa-tasking, not every single retry.

Notes are for non-blocking context on a complete, valid task, for example a label that did not load. Notes never make up for a missing reply, a missing score or an unfinished comparison.

## 9. Check your understanding

Question: What decides the preference when both arms did the work about equally?

Answer: Naturalness for your region: which reply sounds like a real local person writing to this persona. The preference is usually slight in that case.

(Wrong options: the longer reply, even when both completed the same work; whichever reply uses the most formal language.)

## 10. Pay and logistics

- Pay: hourly. Your rate is in your project offer.
- How time is tracked: by keeping your claimed Stoa task open on the platform while you work.
- Claim first: claim the task before you open Hatch. Time in Hatch without an open task is not tracked or paid.
- Payments follow the platform's standard payment process.
- Paid time per task: 20 minutes, covering the whole entry (Hatch plus the form).